

# Operating systems II

## CS 2506

Dr. Dan Grigoras

School of Computer Science and  
Information Technology

# Computing systems

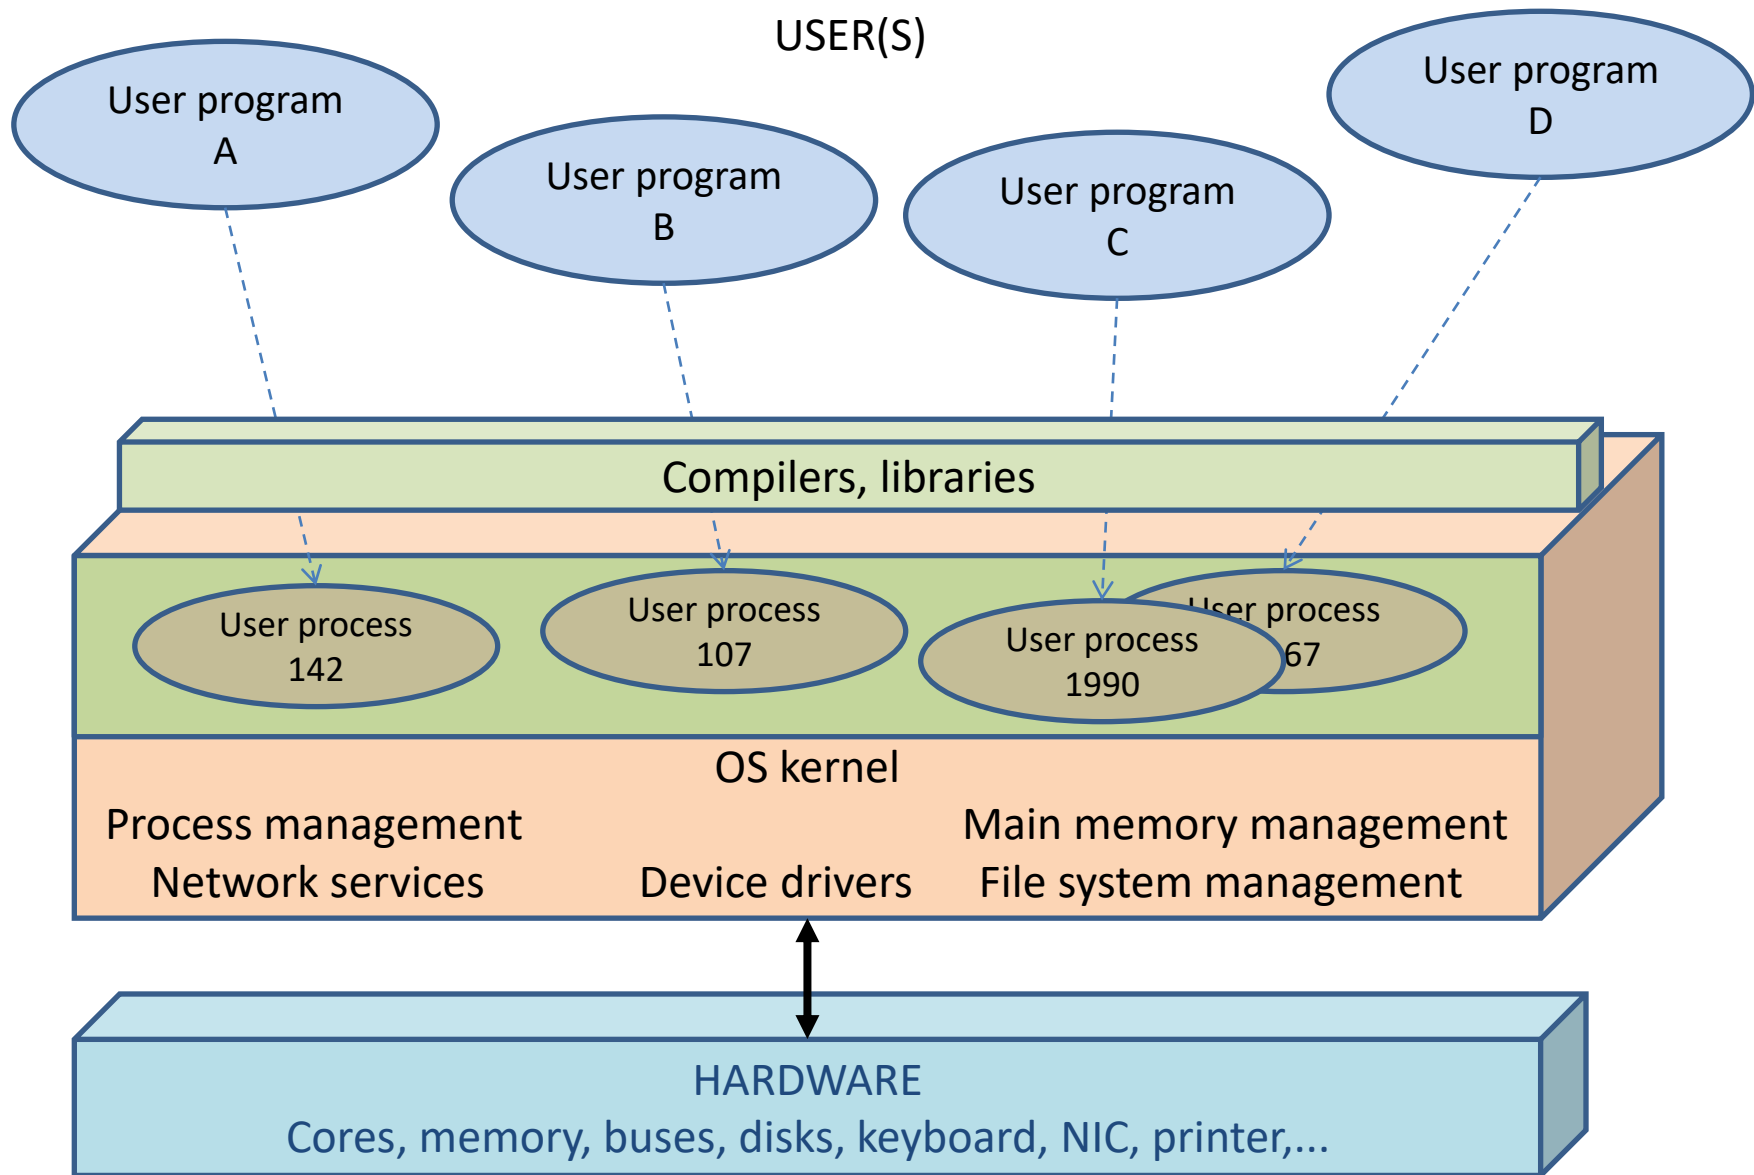
- What is a computer now? Hardware, software, but connectivity plays an ever-increasing role by giving access to remote resources.
- The computer architecture and configuration changed during time:
  - CPU/processor/core → many cores, homogeneous or heterogeneous;
  - Main memory (plain) → memory hierarchy (cache, main memory);
  - Storage (local or remote), disk → array of large disks, SSD;
  - Bus → multiple, faster buses;
  - Peripherals → + Sensors
  - Stationary or mobile
- ...or a *cloud instance* provided as a service.
- *Question*: how are the hardware resources allocated to applications and managed?

# Allocation of computing resources

- The OS manages applications as *processes*. The OS purpose is to allocate the necessary resources for execution and to manage each process during its life cycle.
- Regarding resources, generally, there is more than one physical unit of a type of resource in a computer configuration.  
*Examples:* many cores, memory pages, disk sectors, networking cards, cameras.
- Resource allocation to processes is carried out in terms of units and is managed by specific OS services during processes life-time.
- Many operating systems are using virtualisation of resources such as cores and memory pages for more efficiency.

# Computing resources management

- OS goals:
  - To meet user requirements
    - least execution time, quick response;
    - least cost;
    - fairness;
    - best possible user experience;
    - security.
  - To meet system admin requirements
    - optimal use of resources;
    - least energy use (green computing);
    - maximise revenue.
- The system admin can be the user/owner but not necessarily in all cases – see rack-scale and cloud computing.



# Computing system layers

- The bottom layer is the hardware; it accepts basic commands such as “seek the disk arm to track 79, select head 3 and read sector 5”. Software that interacts directly with the hardware is non-portable and contains low level details: data for control and state registers, interrupt priorities, DMA starting address,...
- The OS kernel has several key components:
  - *Device drivers* are the hardware units’ managers; they hide the low level’s details.
  - *Process management* handles processes, allocating them resources and scheduling their execution.
  - *Memory management* is responsible for physical memory and virtual memory management.
  - *File sys manager* is the code that organizes the data storage into files and directories, hiding low level details re disk blocks, for example.
  - *Network services* provide host-to-host and process-to-process communication across network.

# Access to kernel services

- The kernel is a collection of services, some of which being accessible to user processes. They offer functionality and a higher level of abstraction of the computer. The repertoire of commands supported by the kernel defines the “virtual machine” which is platform-independent.
- To enter the kernel, the user process makes a *system call by executing a trap instruction of some sort*.
- This instruction switches the processor into a privileged operating mode (*kernel mode*) and *jumps into the kernel through a well-defined trap address*.
- Parameters passed with the trap instruction tell the kernel what service is requested.
- When the function is completed, the processor flips back to its normal operating mode (*user mode*) and *returns control to the user process*.

# Trap instructions

- There are functions that require specific knowledge of handling resources – control registers, state register, sequence of operations, and a certain degree of protection – resources shared by several users/processes.
- These functions are coded as service routines; they are also known as *system calls*.
- The sequence of steps for a system call is:
  - The system call is invoked by the user process;
  - OS function is performed;
  - Control returns to the user process.
- Trap instructions are used to implement system calls.



# Library functions

- User programs have access to libraries and can include in their code functions of these libraries – linked in the executables.
- Some library functions use system calls. The function itself parcels up the parameters correctly and then performs the trap.
- The function works as a wrapper for the system call.
- Example:

`printf(“hello world”); ➔ write(1, “hello world”, 11);`

# OS classes

- There isn't one OS that fits all computing devices!
- The large range of computing devices requires customized OS.
- *General purpose computers* run general purpose OS: Unix, Linux, Windows,... that deal with different kind of processes, such as interactive, compute-bound or transaction-based processes.
- *Mobile devices*, such as smart phones or sensors, run OS such as iPhone OS, Android, TinyOS, concerned with power saving and user experience. Their processes are event-driven.
- *Application-oriented OS*: Robot OS, Home OS, Internetware
- *Embedded systems* (IoT) run scaled down versions of OS, real-time, event-driven.

# Trends

- Rapid advances in hardware:
  - peripheral devices are integrated onto a die, each with complex, varying programming models, that execute specialized firmware with little OS integration.
  - for energy saving, cores, devices, and memory can be powered on/off at any time during the system operation.
  - large, non-volatile main memories will likely become prevalent and might have greater capacity than today's disks and disk arrays.
- Applications are changing:
  - many applications use cloud (data center) services (or microservices), individual users' mobile devices, sensors and actuators of the Internet of Things. Do we need an Internet OS?
  - Enterprise applications such as file servers, relational databases, and big data analytics now come pre-packaged in a *rack-scale software appliances*, delivered to customers who have only to plug in and use the system.
  - Clouds are running applications in containers or VMs with their own OS.

# Suggested readings

1. H. Mei, Y. Gao: *Operating Systems for Internetware: Challenges and Future Directions*, IEEE 38th International Conference on Distributed Computing Systems (ICDCS) 2018  
<https://ieeexplore.ieee.org/document/8416401>
  2. D. Milojicic and T. Roscoe: *Will OSs in 2025 still resemble the Unix-like consensus of today, or will a very different design achieve widespread adoption?*, *Computer*, Volume: 49, Jan. 2016, pp. 43-51,  
<http://ieeexplore.ieee.org.ucc.idm.oclc.org/document/7383138>
- My comments: good overview of existing operating systems and technology advances that will have a great impact on operating systems architecture and operation. These papers can suggest interesting avenues of further study and research.

# Course goals and methodology

- Goals
  - understand and learn the main services of an operating system:
    - processes and threads management, including scheduling and load balancing, IPC;
    - cores, memory management;
    - device drivers;
    - file management system.
  - understand and learn features of different models of operating systems:
    - general purpose;
    - mobile;
    - sensor.
- Methodology
  - attending the lectures and engaging in discussions on lectures' topics.
  - carrying out the lab work.
  - using recommended references to learn more about specific topics.
- Most of the topics are comprehensively presented in the book by A. Tanenbaum and H. Bos: Modern Operating Systems, Pearson, 4<sup>th</sup> edition, 2014.

# Canvas CS 2506

- Grading
  - Continuous assessment: 20%
  - Written exam: 90 min 80%
- Lectures: 50 min + 5 min review & questions
- Office: G65, Western Gateway Bldg
- Email: grigoras@cs.ucc.ie, Subject: CS 2506
- Office hours: by appointment

# Plagiarism

1. Plagiarism is presenting someone else's work as your own. It is a violation of UCC Policy and there are strict and severe penalties.
2. You must read and comply with the UCC Policy on Plagiarism [www.ucc.ie/en/exams/procedures-regulations/](http://www.ucc.ie/en/exams/procedures-regulations/)
3. The Policy applies to *all* work submitted, including software.
4. You can expect that your work will be checked for evidence of plagiarism or collusion.
5. In some circumstances it may be acceptable to reuse a small amount of work by others, but *only* if you provide explicit acknowledgement and justification.
6. If in doubt ask your module lecturer *prior* to submission. Better safe than sorry!