
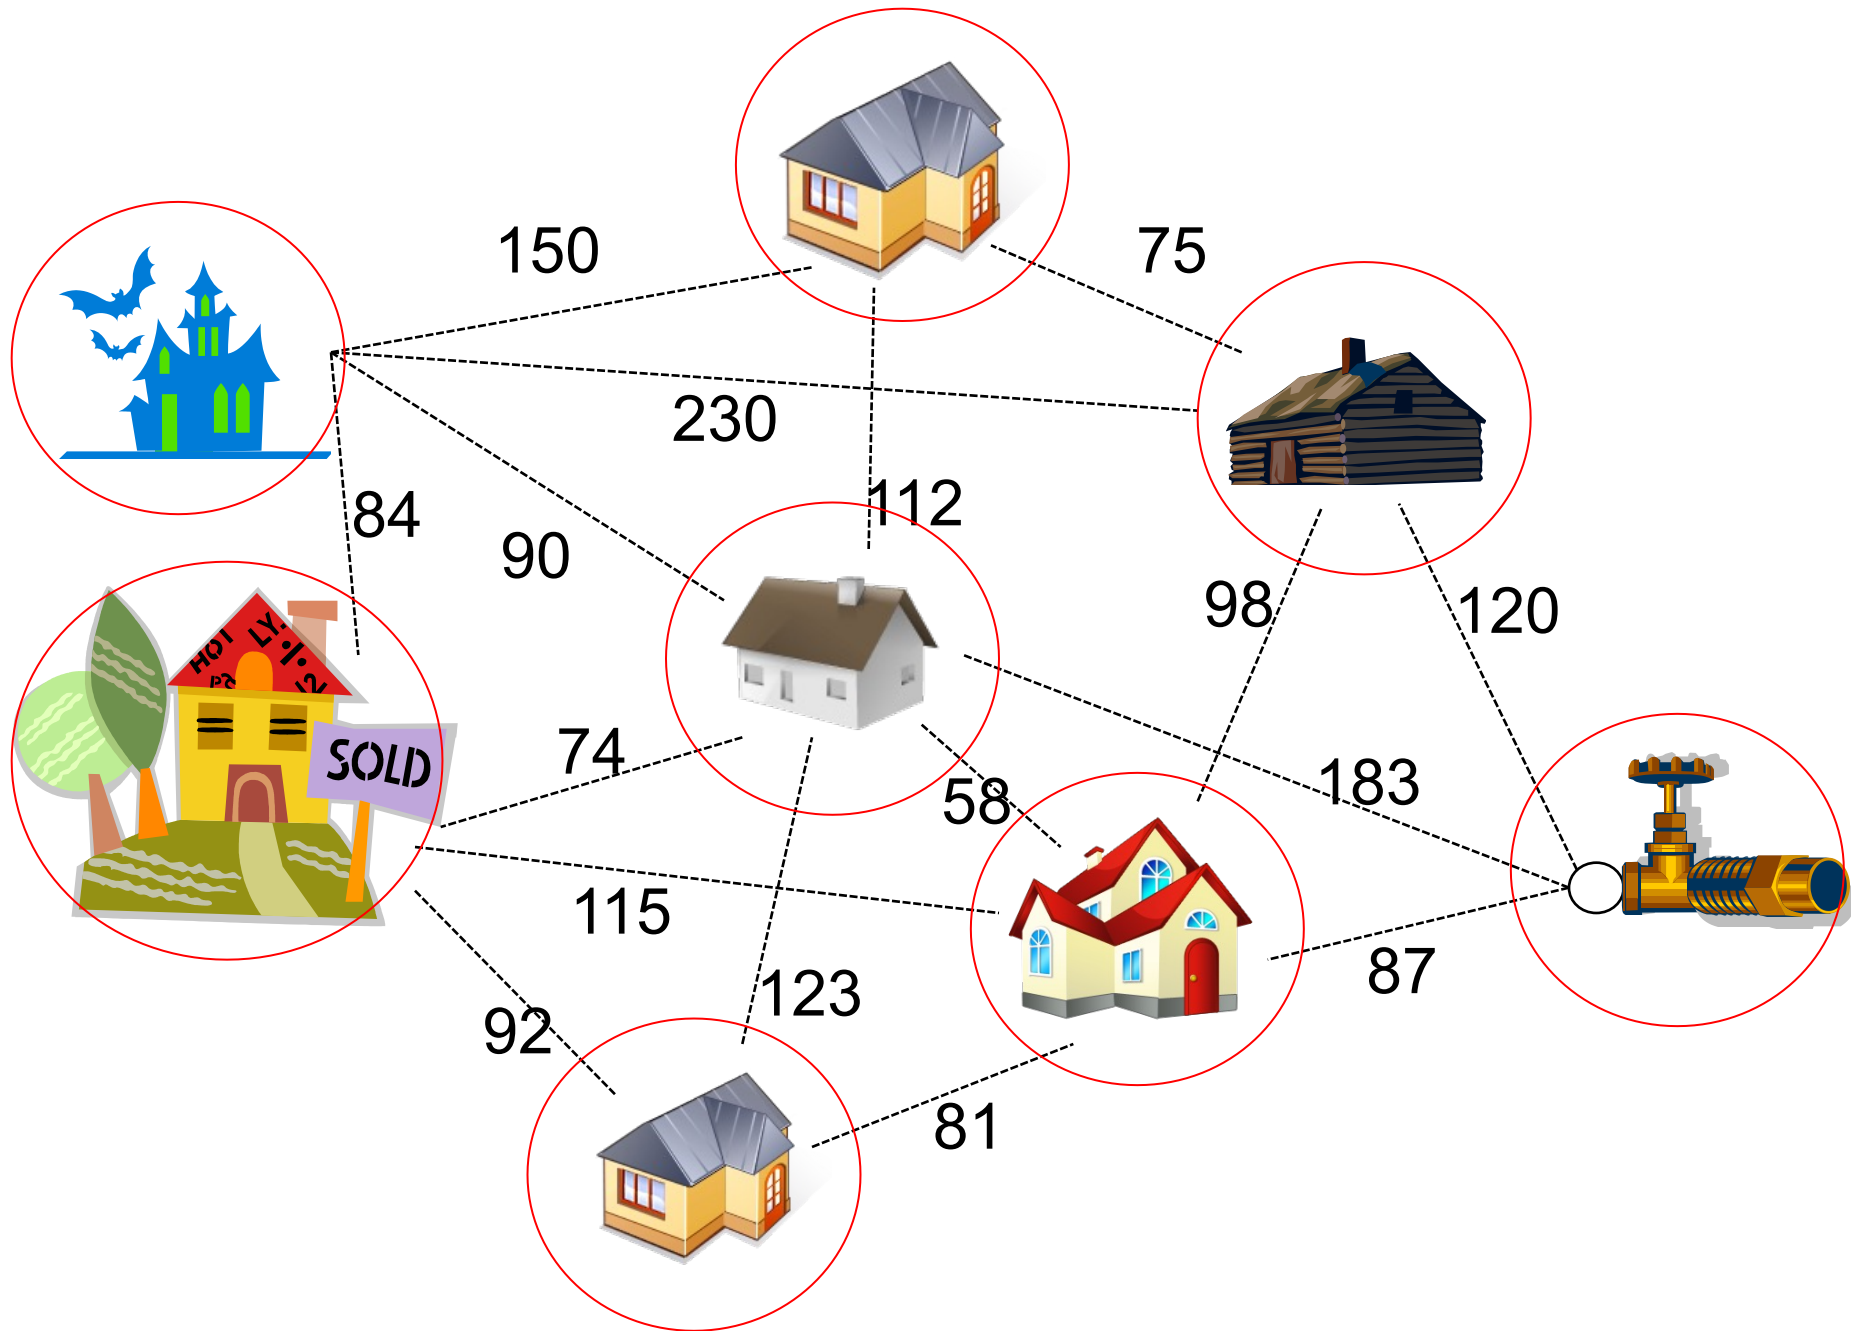


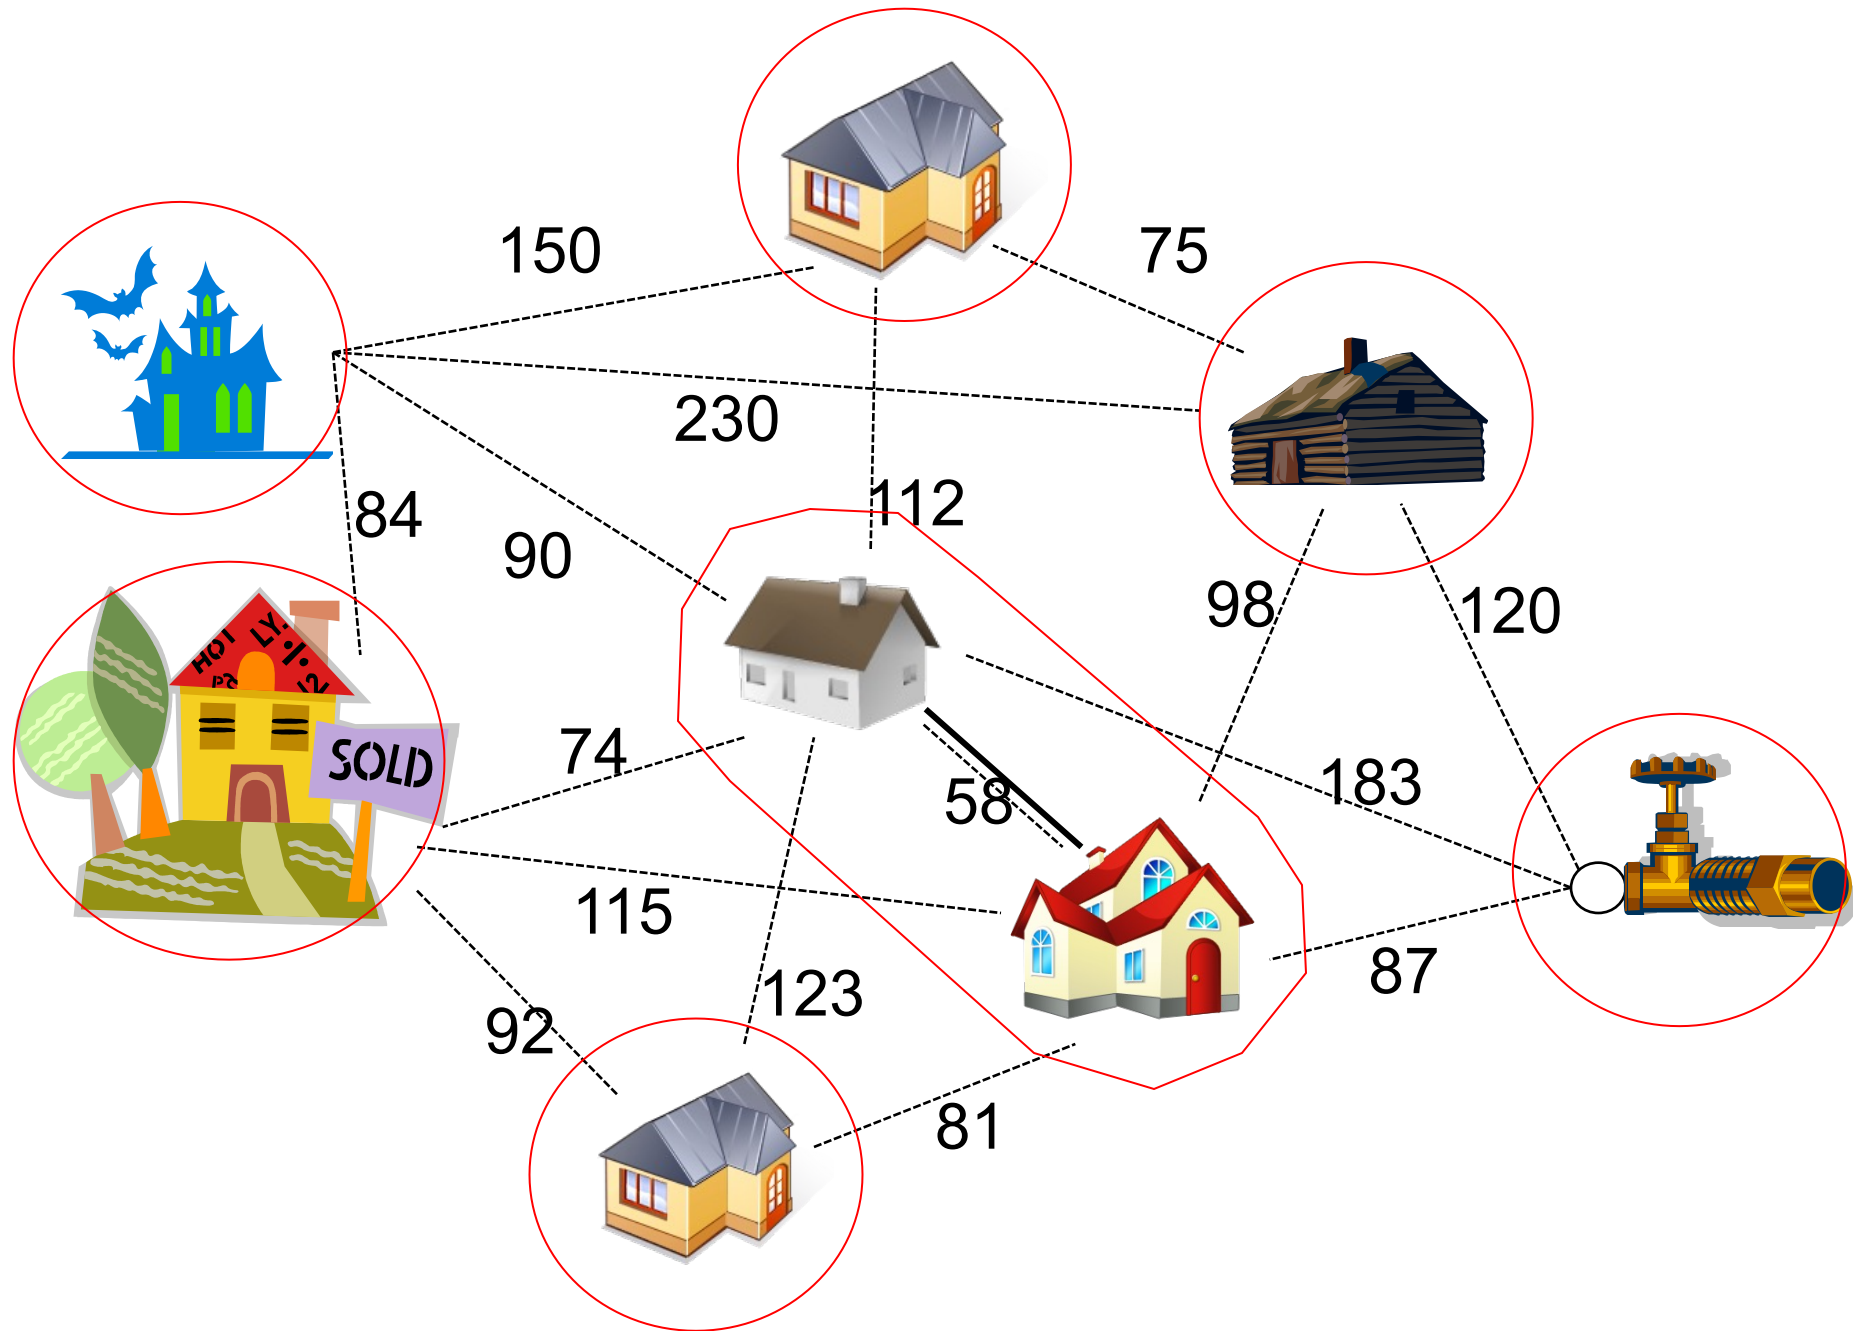
Minimum Spanning Trees: Kruskal's Algorithm

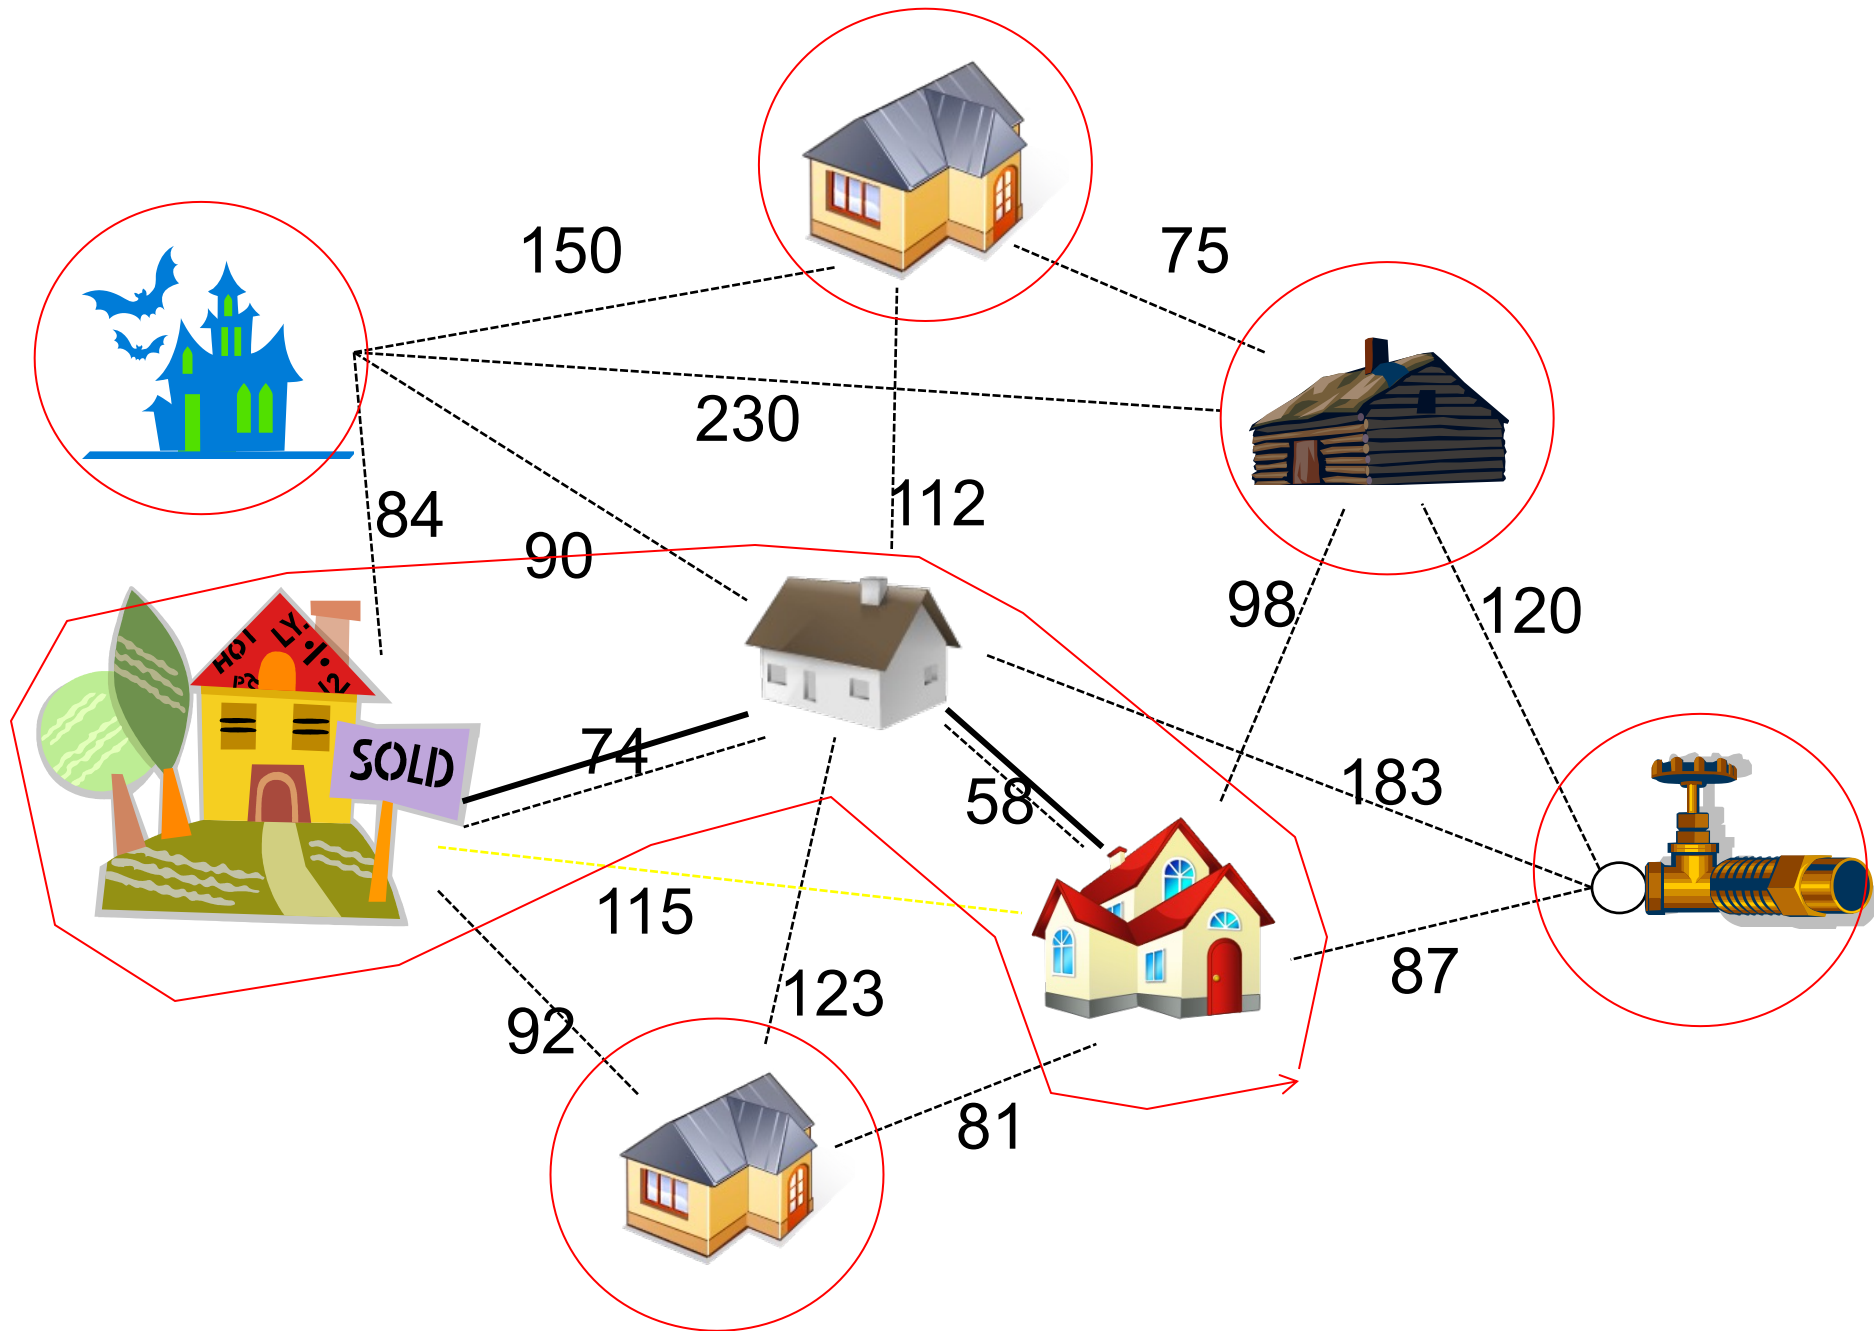


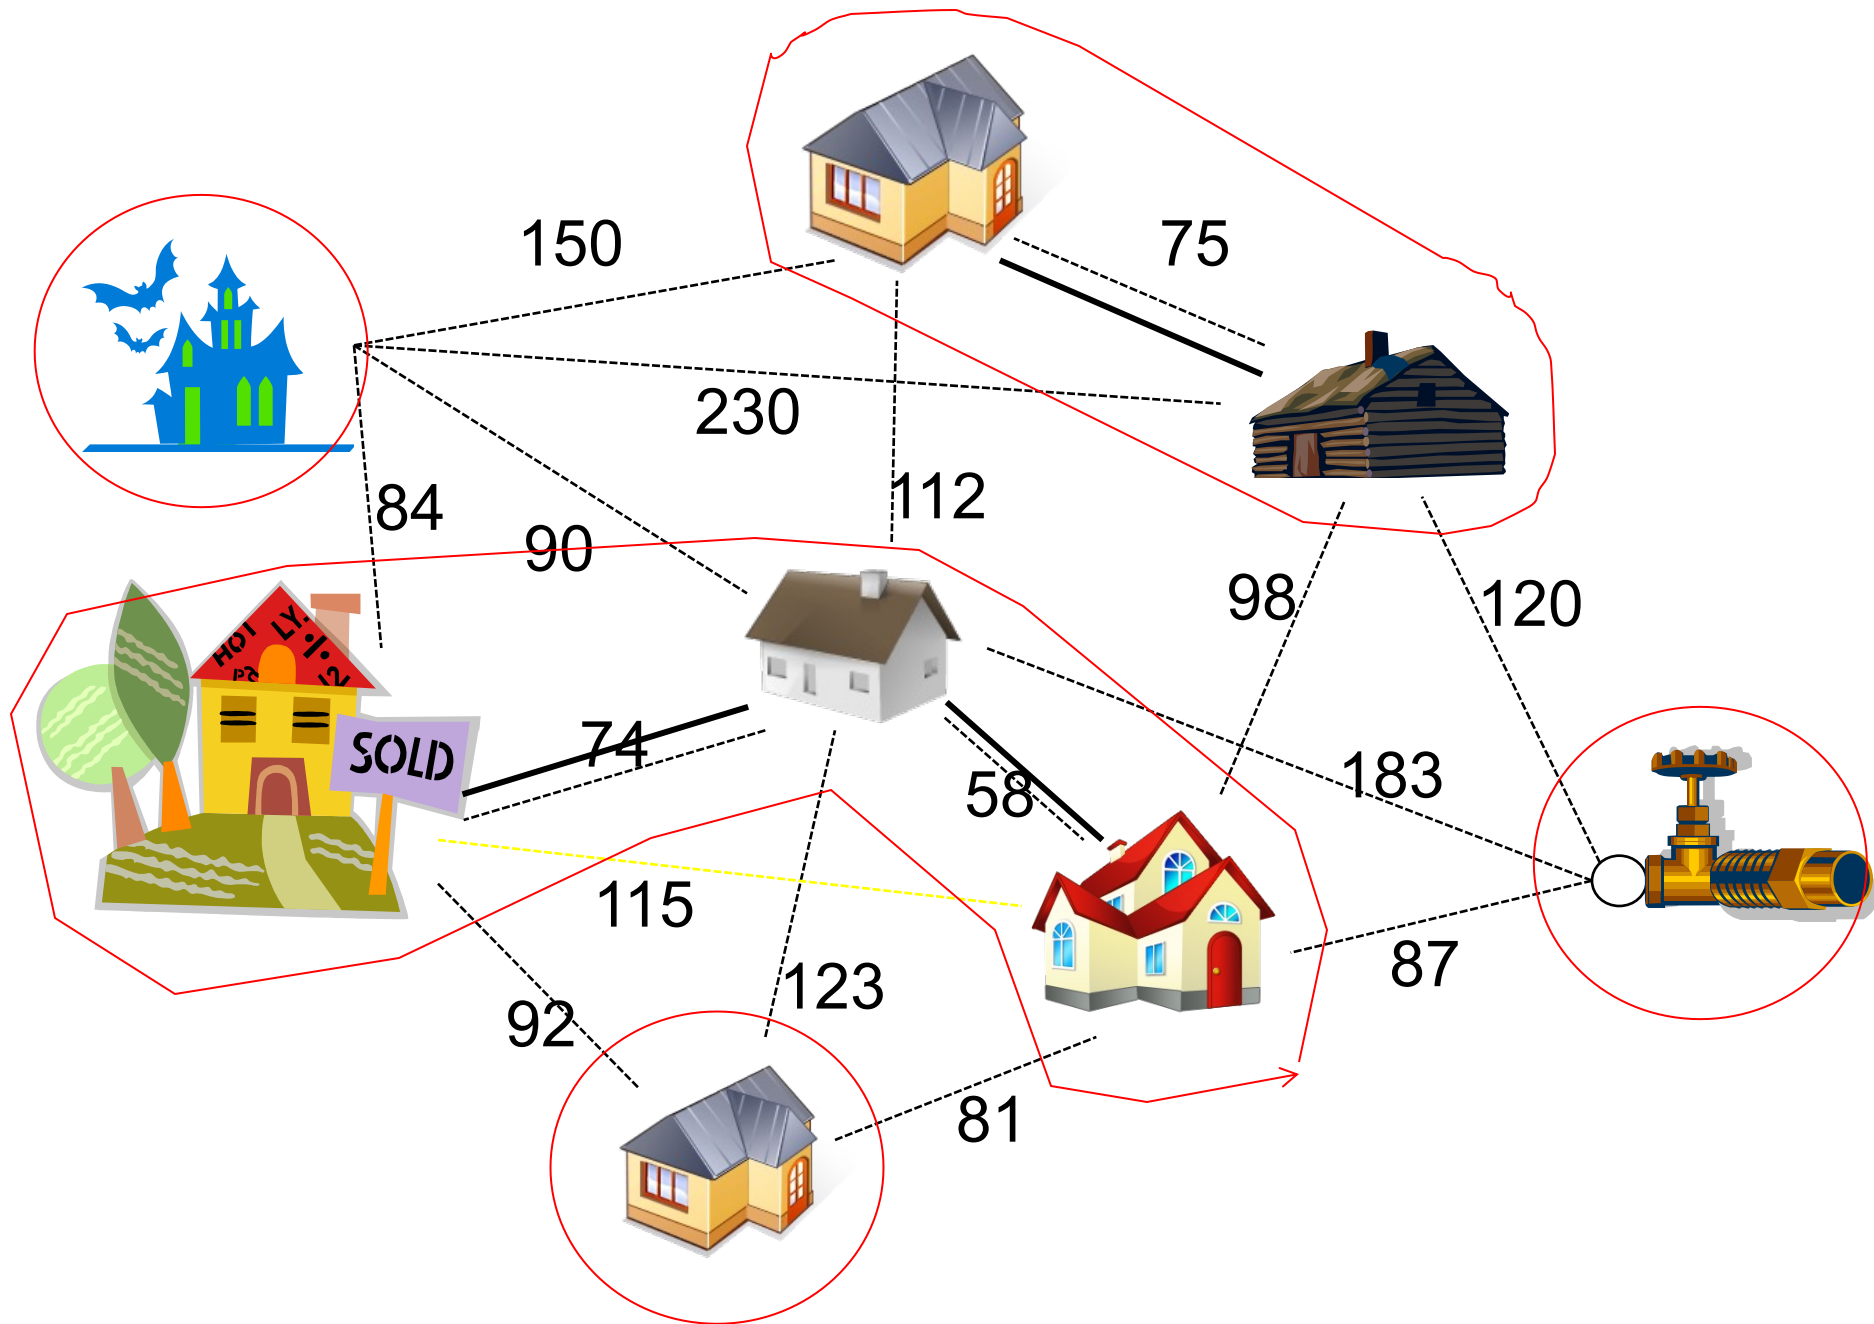
Prim's algorithm worked by gradually building a single tree one edge at a time until all vertices in the graph are in the tree.

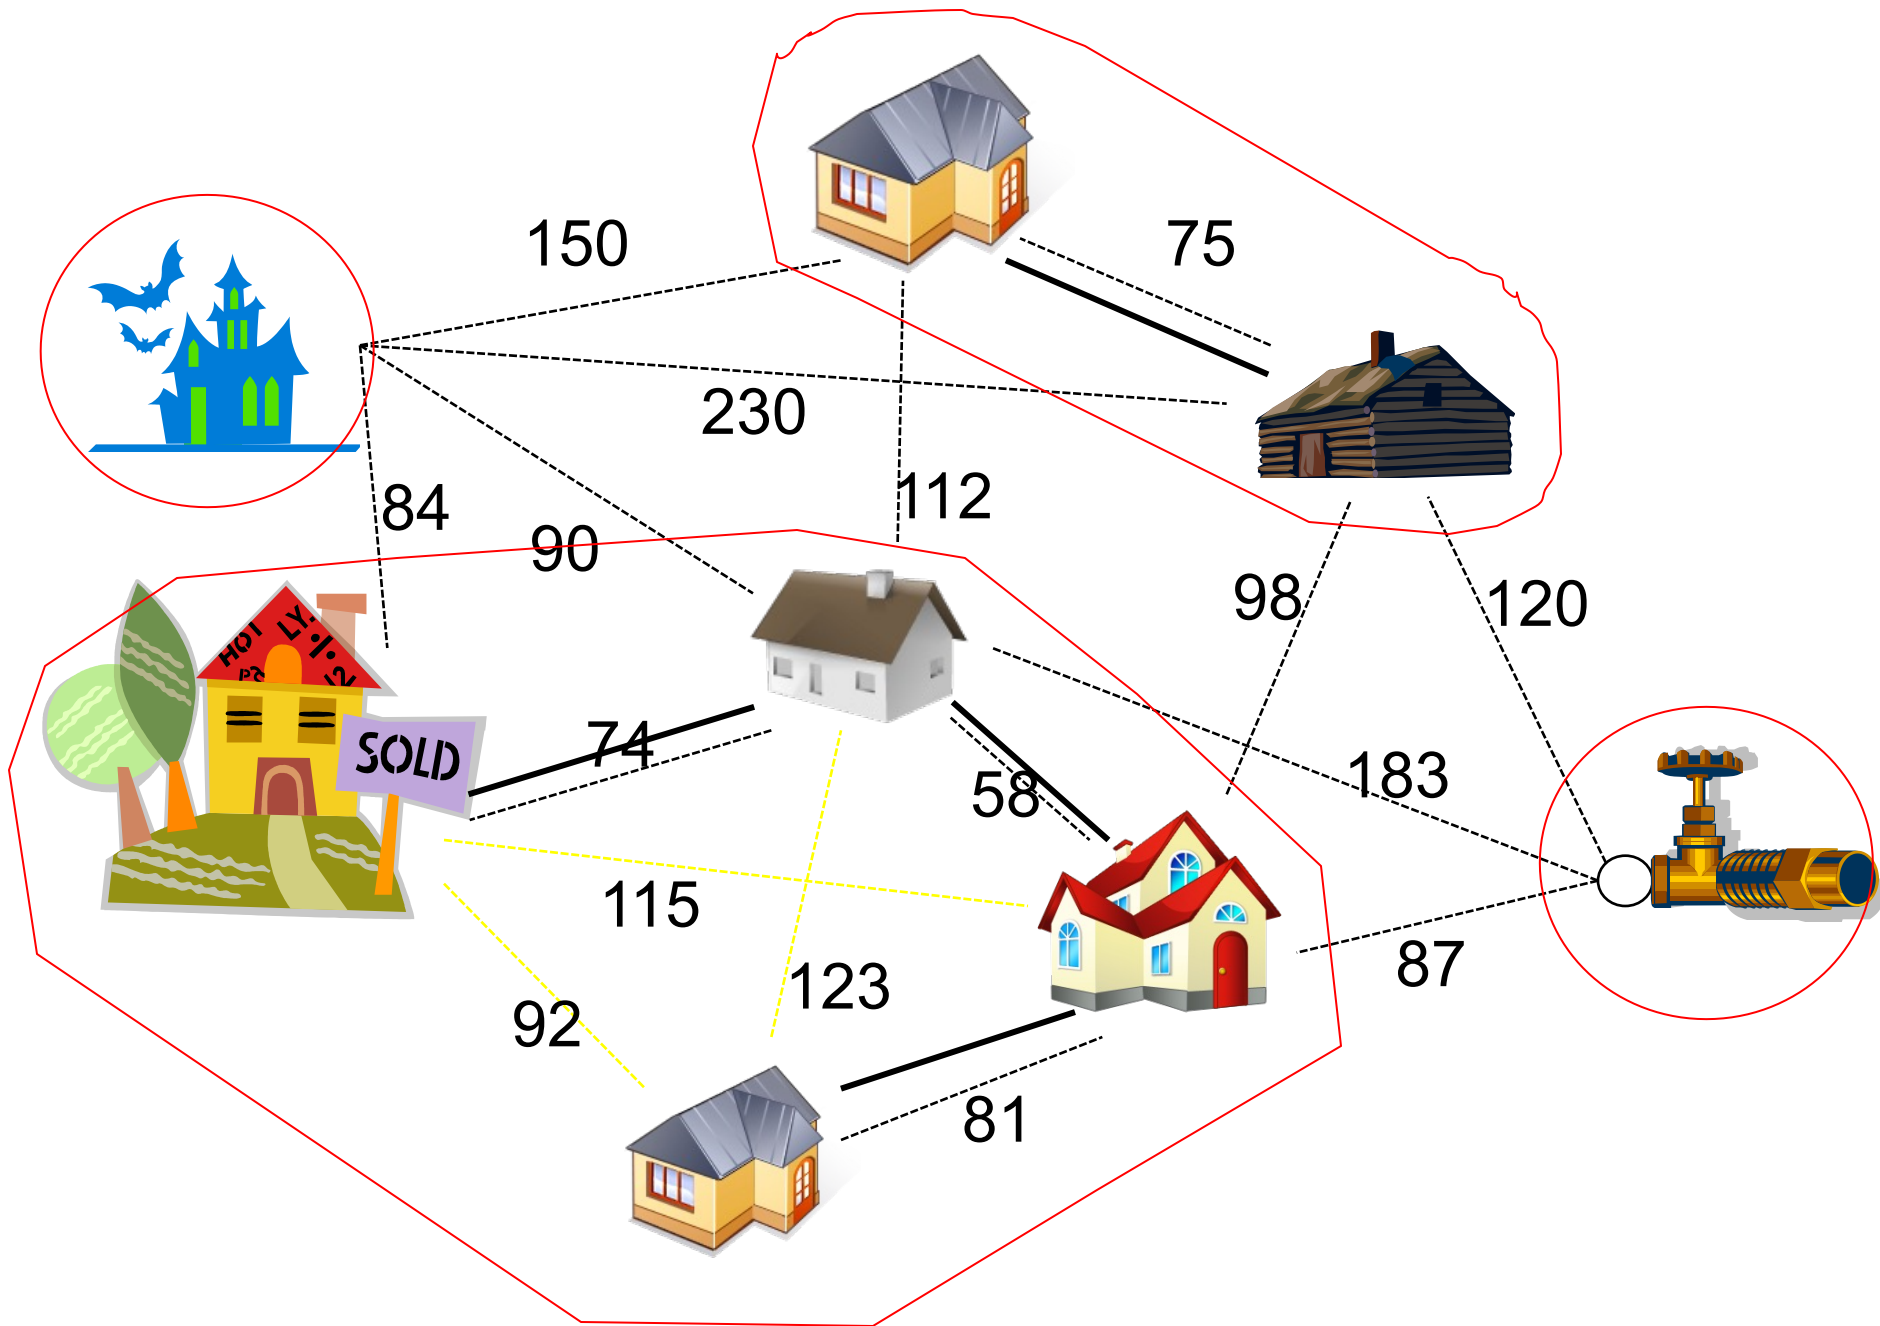
There is another approach, Kruskal's algorithm, which works by joining trees together until all vertices from the graph are in a single tree.

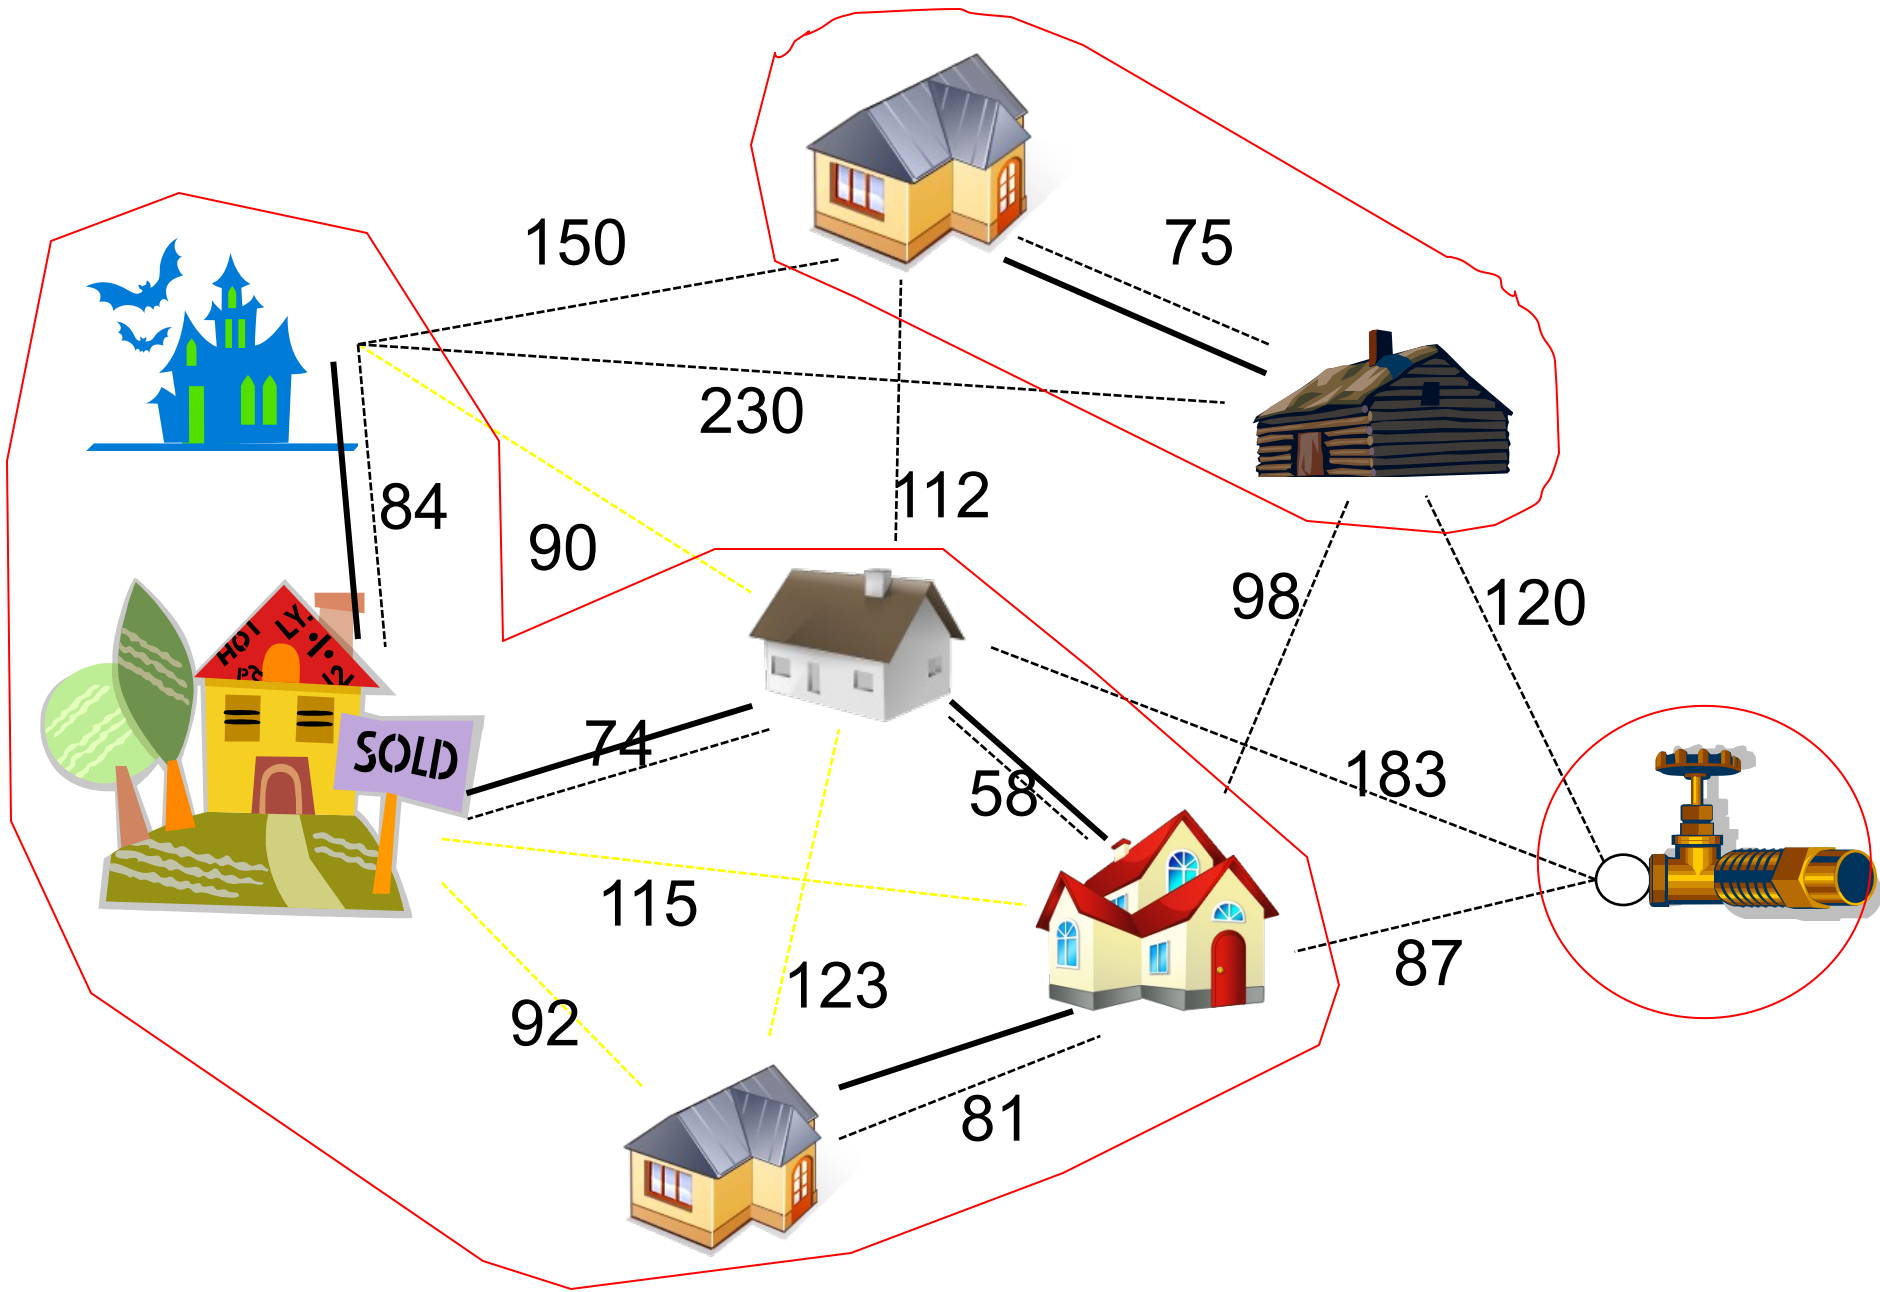


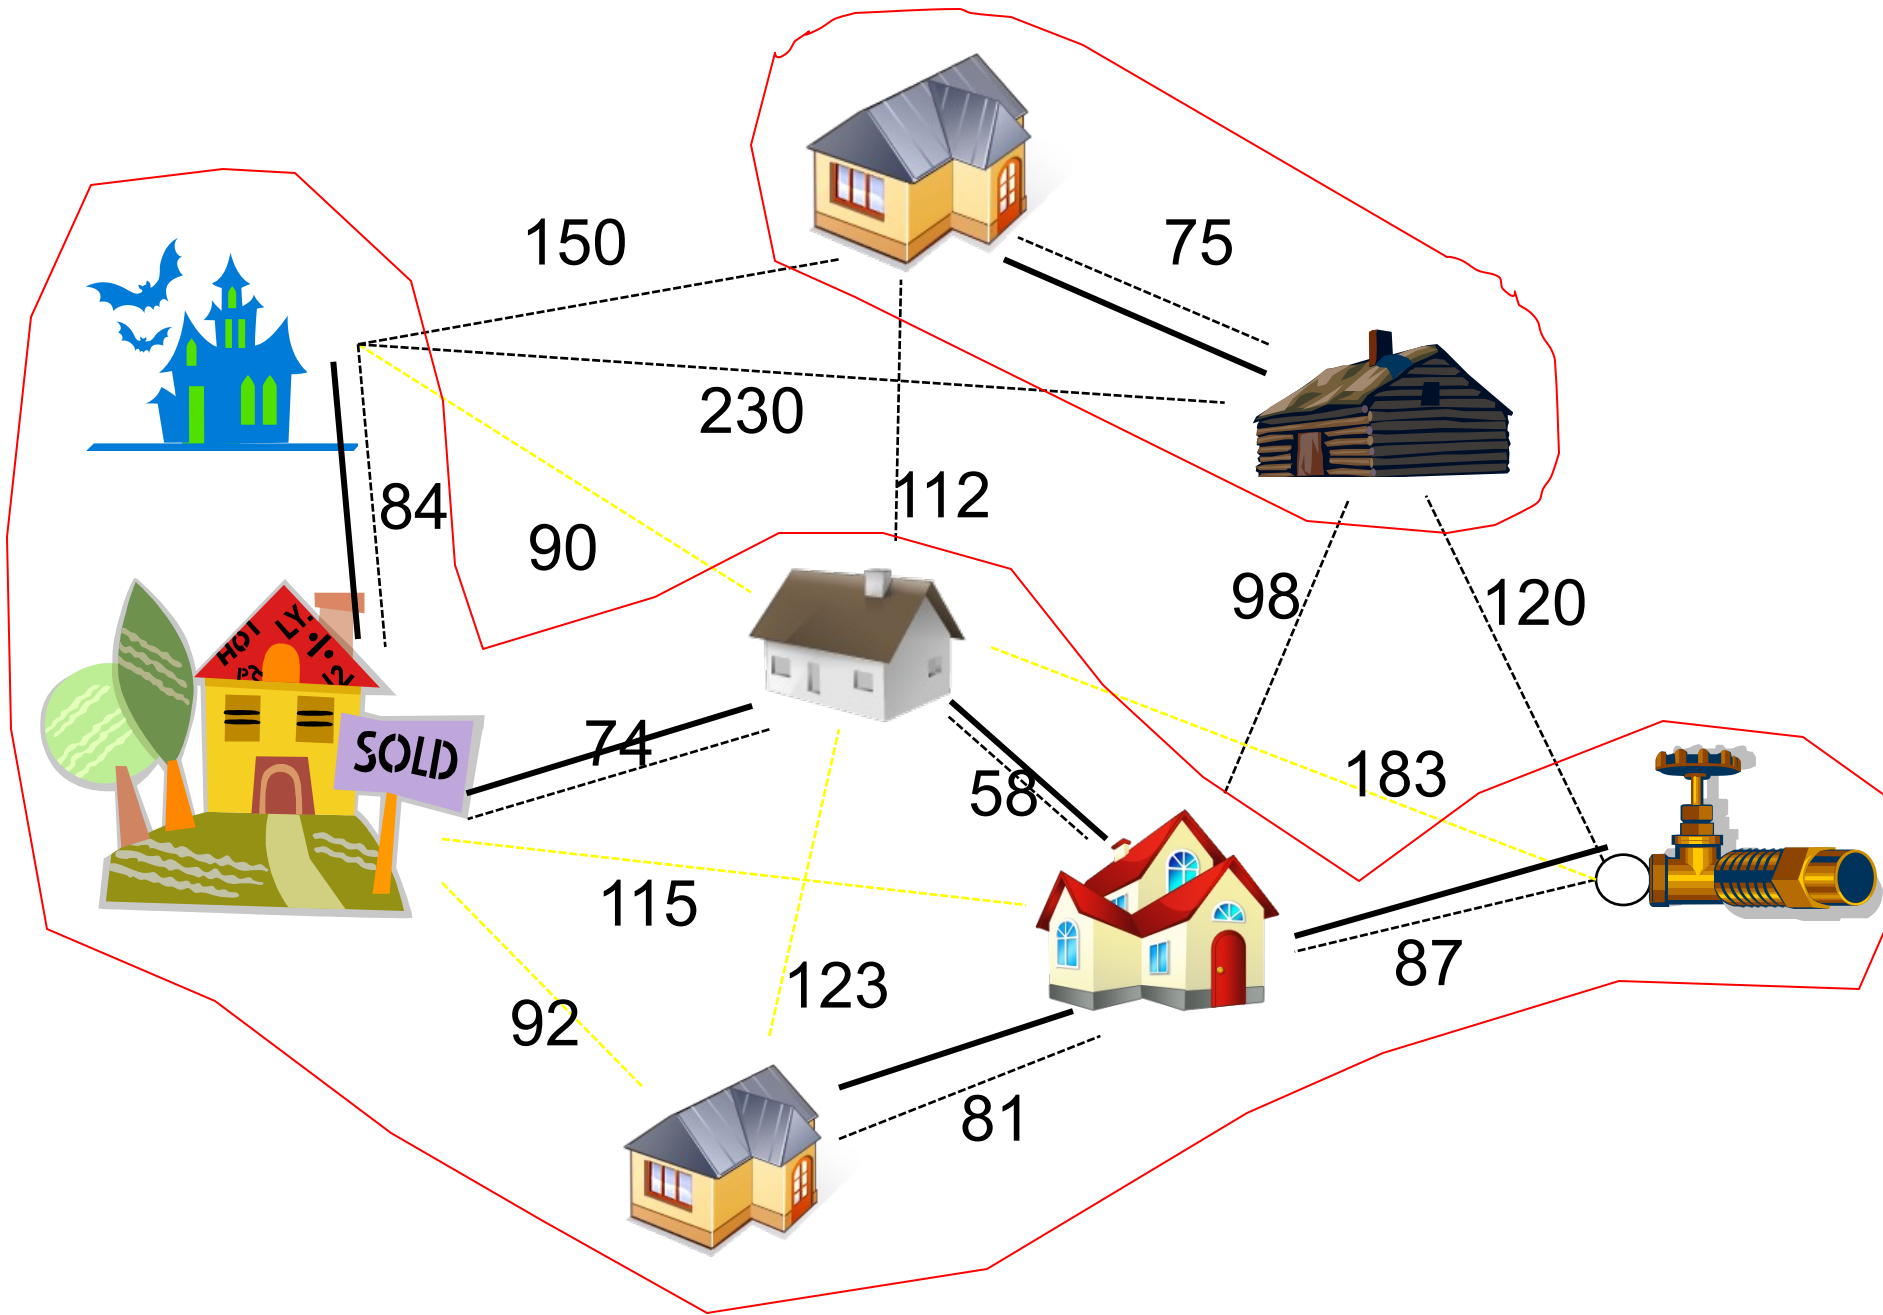


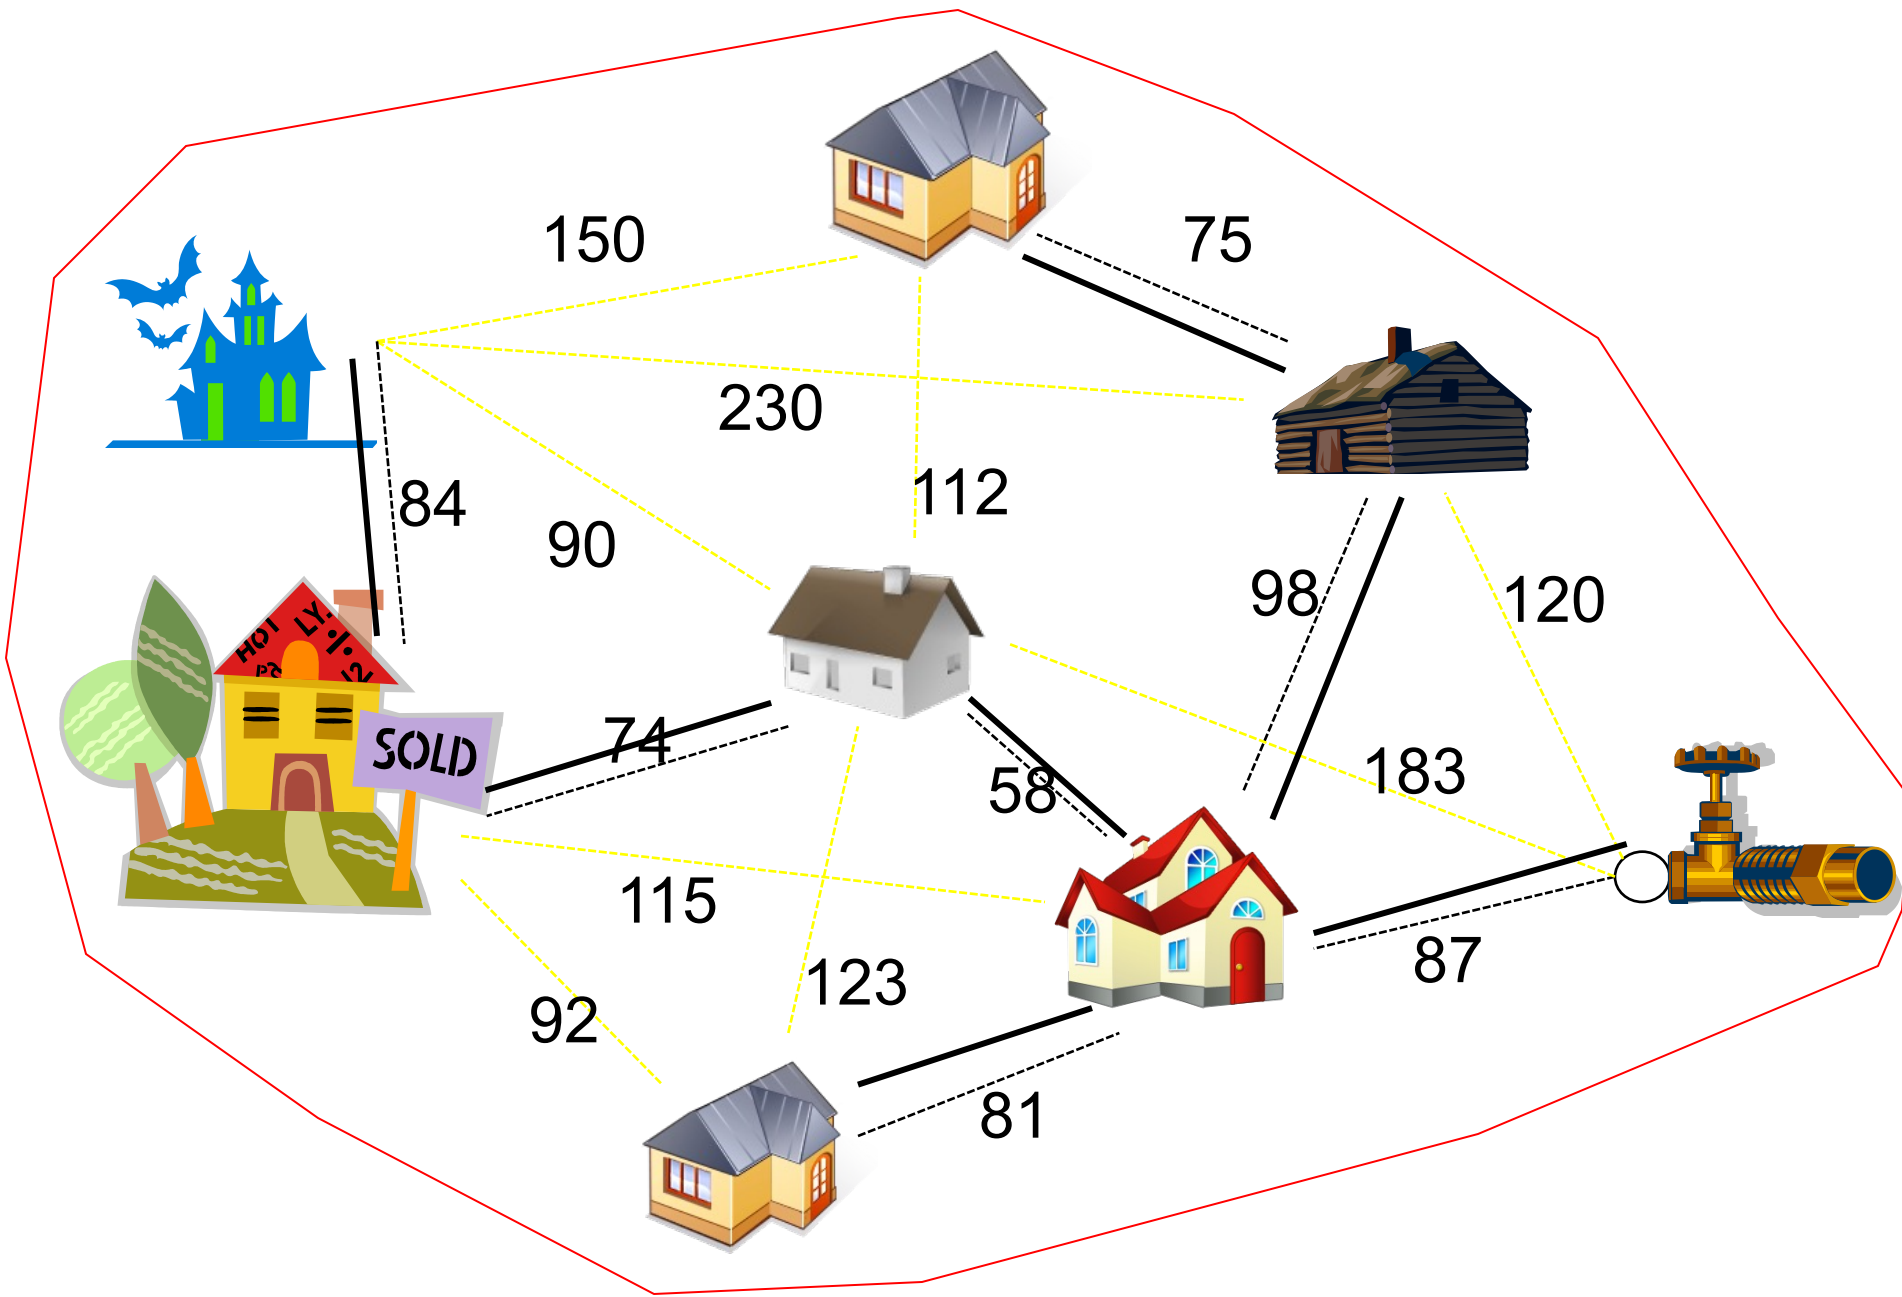











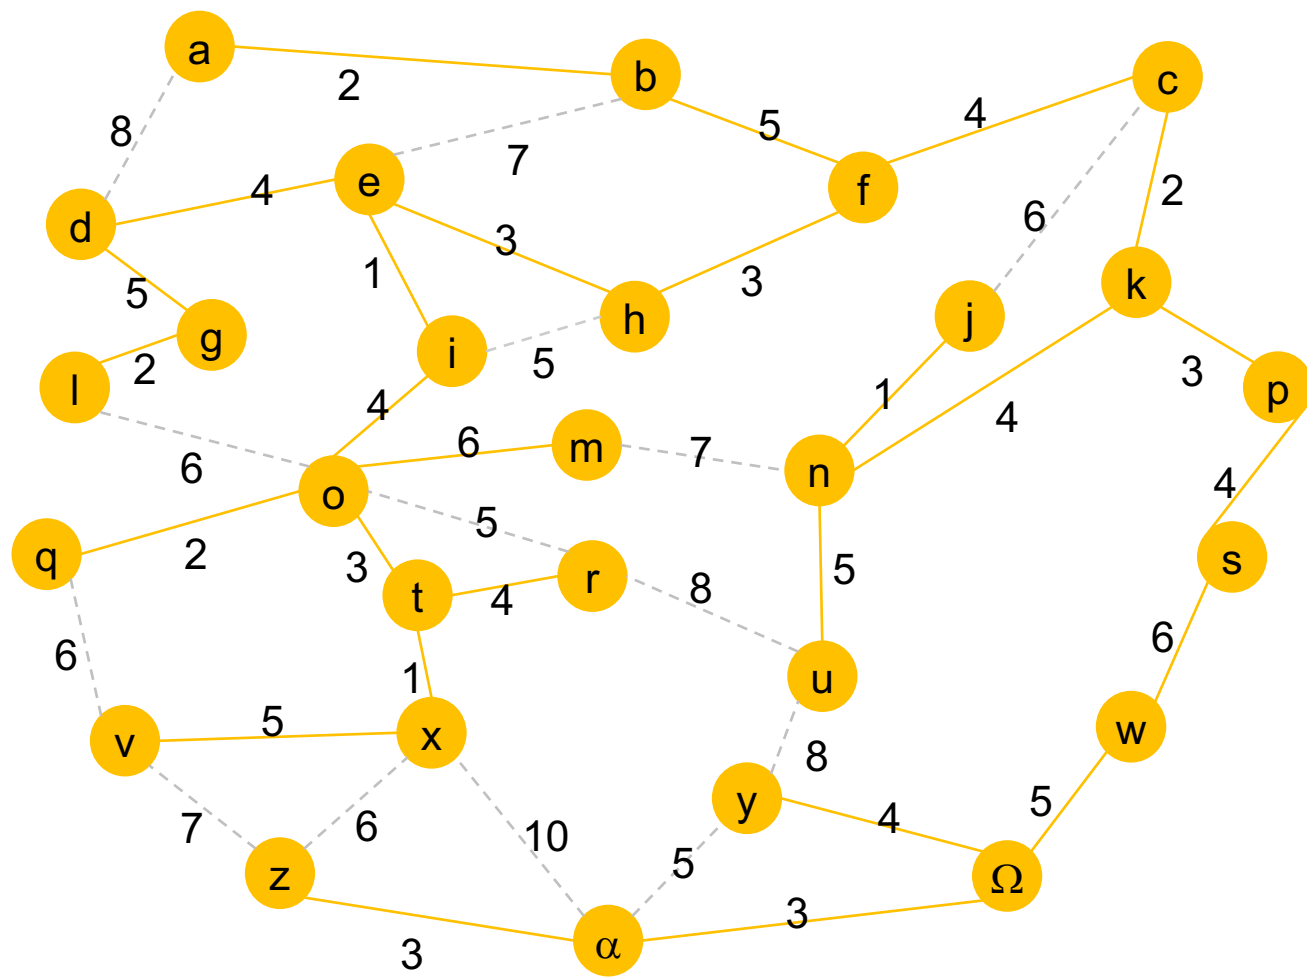


```
kruskal(): #pseudocode version 1
create an empty mst //a list of edges
for each vertex in the graph
    create a separate tree
for each edge in the graph in increasing order of cost
    if edge joins two separate trees
        add edge to the mst
        merge the two trees into one tree
return mst
```

Sketch of proof of correctness

- if the graph is connected, 'mst' contains every vertex, so spans the graph
- 'mst' is a tree (never adds an edge that creates a cycle)
- proof that it has minimum cost is similar to that for Prim's algorithm





```
kruskal(): #pseudocode version 1
create an empty mst
for each vertex in the graph
    create a separate tree
for each edge in the graph in increasing order of cost
    if edge joins two separate trees
        add edge to the mst
        merge the two trees into one tree
return mst
```

How do we implement this efficiently?

kruskal(): #pseudocode version 2

create an empty list *mst* // to contain edges in final MST

create an empty *dictionary* // for each vertex, state its current tree

create a PriorityQueue *pq* //to order graph edges by cost

for each edge *e* in *G*

add ($w(e)$, *e*) into *pq*

for each vertex *v* in *G*

create a tree *t* for *v* and add (*v*:*t*) to *dictionary*

while $\text{length}(mst) < |V|-1$ and *pq* is not empty

remove *e*, the minimum cost edge from *pq*

get the trees t_1 , t_2 for *e*'s vertices from *dictionary*

if t_1 and t_2 are different

add *e* into *mst*

join t_1 and t_2 into a single tree *t* and update *dictionary*

return *mst*

How do we implement this efficiently?

Representing the trees

Implement each tree as a stack (or any sequence with $O(1)$ updates).

To join two trees, find the smaller one, then until that stack is empty:

- pop a vertex

- update its dictionary entry to point to the bigger tree

- push onto the bigger tree.


```

def kruskal(self):
    mst = []
    n = self.num_vertices()
    whichtree = {}

    for v in self.vertices():
        vtree = Stack()
        vtree.push(v)
        whichtree[v] = vtree
    pq = PQHeap()

    for e in self.edges():
        pq.add(e.element(), e)
    while len(mst) < n and pq.length() > 0:
        key, e = pq.remove_min()
        (x,y) = e.vertices()
        xtree = whichtree[x]
        ytree = whichtree[y]
        if xtree != ytree:
            mst.append(e)
            self._jointrees(xtree, ytree, whichtree)
    return mst

```

```

def _jointrees(self, xtree, ytree, whichtree):
    if xtree.length() < ytree.length():
        target = ytree
        deltree = xtree
    else:
        target = xtree
        deltree = ytree
    while deltree.length() > 0:
        v = deltree.pop()
        whichtree[v] = target
        target.push(v)
    del deltree

```

Complexity:

creating the dictionary: $O(n)$

creating the PQ: $O(m \log m)$ - the number of edges

at most m times round the loop

for each time around the loop, $O(\log m)$ to remove the min-cost edge,
 $O(1)$ to get the trees for the two vertices and test if they are different, and
then the cost of merging the trees.

So $O(m \log m)$ overall to handle the PQ, ignoring the tree merge cost.

Each time two trees get merged, the vertices in the smaller tree move. Each individual vertex move is $O(1)$ – `pop()`, dictionary update, `push()`. Since we only move vertices from the smaller tree, each vertex that is moved moves to a tree that becomes at least double the size. There are n vertices, so each vertex can change trees at most $O(\log n)$ times (since $1 \cdot 2 \cdot 2 \cdot \dots \cdot 2$ for $\log n$ times is n). So amortised cost of tree merging is $O(n \log n)$.

So in total $O(n \log n + m \log m)$

But m is $O(n^2)$, so $O(\log m) = O(\log n^2) = O(2 \log n) = O(\log n)$

So algorithm is $O((n+m) \log n)$

doesn't matter
if it is sparse or
dense ...

Next lecture

Further graph algorithms